

Country-Scale Offline Speed-Limit Inference on Smartphones: A Runtime Data Architecture Study

Anonymous Authors

Submission prepared for double-blind review

Abstract—This paper evaluates runtime data architectures for low-latency, offline-capable smartphone speed-limit inference using country-scale OpenStreetMap (OSM) data. We identify four data processing scenarios that are compared to find an optimal query and update scenario along three typical query conditions. Hence, we identify the preferred operating point for practical deployment for on-device speed-limit inference as required for Intelligent speed assistance (ISA) applications that are mandatory for new vehicles in the EU since 2022 and open the way for a smartphone and open data based retrofitting approach for vehicles that are already in traffic. Thereby we can strongly support the EU traffic safety vision of reduced crash severity.

Index Terms—Intelligent transportation systems, intelligent speed assistance, mobile GIS, OpenStreetMap, spatial indexing, offline-first

I. INTRODUCTION

Excessive speed remains a major crash-severity factor, and the deployment of Intelligent Speed Assistance (ISA) is now required in new EU vehicle types [1]–[3]. However, the installed vehicle base is substantially larger than annual new registrations, which motivates retrofit-capable software solutions that can run on smartphones with open geospatial data.

Most map-matching research focuses on inference models under sparse/noisy trajectories. In practical mobile deployment, a separate constraint is often dominant: country-scale map data must be queried with low latency under limited memory, storage I/O, and intermittent connectivity. This paper targets that runtime architecture problem.

The practical motivation is scale. EU passenger-car stock remains above 260 million vehicles, whereas annual registrations are only a fraction of that stock [4], [5]. In parallel, smartphone internet usage is high in both urban and rural contexts [6], making mobile offline-first assistance a feasible deployment channel for legacy vehicles.

The study is positioned for ITS tracks focused on infrastructure technologies and traffic data systems. The contribution is not a new probabilistic matching model; it is a deployment-oriented comparison of runtime data representations and update behavior for speed-limit inference.

Contributions:

- 1) A controlled comparison of four country-scale runtime architectures (S1–S4) built from identical source data.
- 2) Joint latency evaluation for cloud-like host execution and physical smartphone execution across three query modes.

- 3) A 30-day daily-diff update analysis to compare maintenance effort and update locality.
- 4) A deployment recommendation balancing query latency and incremental update cost.

II. RELATED WORK

Map matching for transport applications has been widely studied, including geometric/topological methods [7] and probabilistic sequence models such as HMM-based approaches [8]–[11]. Later work improved robustness for sparse trajectories and route-level inference [12]–[16].

This paper complements that line of work by isolating the systems layer around map data packaging and spatial access. R-tree indexing [17] is treated as an established mechanism whose deployment characteristics are compared across architecture variants.

Compared with prior ITS map-matching studies, the design focus here is on country-scale runtime packaging and update operability, not only per-query geometric accuracy. This positioning is relevant to production constraints in connected and non-connected operation: an architecture that is fast in a server benchmark may fail under mobile sandbox I/O constraints, while a compact architecture can preserve both latency and maintainability.

III. METHODOLOGY

A. Data scope and regulatory baseline

All scenarios use the same OSM Germany snapshot (`germany-latest.osm.pbf`, 4.70 GB) [18]. Preprocessing retains car-drivable `highway=*` classes and associated speed annotations. Table I combines dataset profile statistics with speed-limit provenance after rule-aware fusion [19], [20].

When explicit speed tags are absent, the legal default baseline for Germany (50 km/h urban, 100 km/h rural) is applied [21]. This rule-aware fallback is necessary because most ways do not carry explicit speed tags.

B. Runtime architecture scenarios

The four scenarios reflect distinct deployment trade-offs. S1 is a global-index baseline that queries country-scale index files directly. Its implementation is simple, but each lookup can trigger broader scans and higher I/O pressure. S2 introduces explicit spatial partitioning with content-addressed tiles. This localizes reads to relevant map partitions and supports selective tile-level refresh when updates arrive. S3 shifts to a single-file embedded database with spatial indexing; query

TABLE I
COMBINED DATASET PROFILE AND SPEED-LIMIT PROVENANCE
(GERMANY SNAPSHOT).

Metric	Value
Input file size	4.70 GB
Car-drivable ways	7,963,481
maxspeed=* ways	2,637,020
zone:maxspeed=* ways	172,116
Explicit speed-limit ways	2,639,693
Area-context objects	129,459
Explicit tagged limit (share)	33.15%
Implicit regulatory tag (share)	0.98%
Default-rule fallback (share)	65.87%

execution performs direct spatial-window retrieval followed by geometric ranking, reducing candidate fan-out and simplifying deployment to one artifact. S4 adds a tile-membership prefilter before spatial-index intersection, aiming to prune dense urban candidate sets at the cost of an additional filtering stage.

For deployment, S2–S4 use immutable file artifacts per build and are therefore suitable for object-storage distribution, stateless worker replication, and autonomous offline execution on mobile devices.

C. Query modes

All scenarios share the same ranking score:

$$\text{score} = d + \lambda_h \cdot \Delta\theta + p_u \quad (1)$$

where d is the geometry distance term, $\Delta\theta$ is bidirectional heading mismatch, λ_h is heading weight, and p_u is an optional unknown-highway penalty.

The three query scenarios are designed to cover practical operating modes for moving-vehicle inference:

- **bbox**: d is computed from haversine distance to the clamped point on each candidate bounding box. This is the fastest approximation and acts as a coarse candidate filter.
- **polyline**: d is computed as minimum point-to-segment distance over full candidate geometry. This provides the highest geometric fidelity, especially on curved and elongated segments.
- **hybrid**: candidates are first ranked by **bbox**; a limited top- N subset is then refined using **polyline**. This aims to preserve near-**polyline** quality while reducing compute cost.

These scenarios capture a practical latency-fidelity spectrum and are evaluated identically for all runtime architectures.

D. Update analysis protocol

Incremental behavior is evaluated over a 30-day daily-diff window (2026-01-22 to 2026-02-20). The analysis measures changed ways, speed-tag changes, and scenario-specific refresh effort proxies (partition invalidation for S1/S2; simulated patch runtime for S3/S4).

The update protocol uses daily replication deltas to approximate operational maintenance effort between full snapshot

TABLE II
MEAN LATENCY MATRIX (ms) AT THE BENCHMARK COORDINATE.

Mode	S1	S2	S3	S4
Online (cloud-like)				
bbox	4815.67	187.28	84.39	63.28
hybrid	8306.77	179.32	27.03	67.77
polyline	8367.55	165.19	35.99	76.33
Offline (device)				
bbox	3970.84	194.49	37.12	757.01
hybrid	2535.44	45.76	23.39	95.96
polyline	2639.76	62.88	44.93	86.48

checkpoints. This allows comparing granular partition invalidation (S1/S2) with patch-like database updates (S3/S4) under the same observation horizon.

IV. EXPERIMENTAL SETUP

Two deployment contexts are measured:

- **Cloud-like host**: MacBook Pro (Apple M4 Max, 48 GB RAM), macOS 15.5.
- **Offline device**: physical iPhone 14 Pro (iOS 26.2.1), data stored in mobile sandbox.

The benchmark uses a fixed Berlin probe point (52.5200, 13.4050), heading 90° , top- $k = 5$, and all three query modes.

Runtime footprint differs substantially by architecture: S1 uses 7 files (5.09 GB), S2 uses 109,333 files (4.61 GB), S3 uses 1 file (2.60 GB), and S4 uses 1 file (3.00 GB).

V. RESULTS

A. Latency comparison

Table II reports mean latencies (ms) across deployment contexts and query modes, while Figure 1 presents the same matrix in scenario-first view.

S3 is the strongest overall device configuration for decision-relevant query modes (hybrid, polyline). Compared with S1, host-side speedups reach $307.3\times$ (S3, hybrid). On device, the ordering remains $S3 < S2 < S4 \ll S1$ for hybrid and polyline.

An important observation is that S4 does not uniformly dominate S3. While S4 applies additional prefiltering intended to reduce dense-area fan-out, the extra stage increases overhead in some device-side conditions, including the **bbox** mode where S4 regresses strongly versus S3. This indicates that multi-stage filtering needs workload-specific tuning and is not automatically beneficial.

To translate latency into a raw location-update budget, the following relation is used:

$$f_{\text{query}} = \frac{1000}{t_{\text{ms}}}, \quad v_{\text{max}} = 3.6 \cdot f_{\text{query}} \cdot \Delta s \quad (2)$$

where t_{ms} is mean latency and $\Delta s = 10\text{ m}$ is the inter-fix spacing assumption.

TABLE III
ON-DEVICE THROUGHPUT ENVELOPE FROM MEASURED LATENCY.

Scenario	Mean ms	Hz	Best $v_{\max}$	Worst $v_{\max}$
S1	3048.68	0.33	14.2	9.1
S2	101.04	9.90	786.7	185.1
S3	35.15	28.45	1539.1	801.2
S4	313.15	3.19	416.3	47.6

TABLE IV
30-DAY DAILY-DIFF UPDATE SUMMARY (2026-01-22 TO 2026-02-20).

Metric	Total	Mean/day
Changed ways (+, −, modified)	1,829,171	60,972
maxspeed-related tag changes	139,999	4,667
Invalidated S1 grid cells	201,355	6,712
Invalidated S2 tiles	80,536	2,685
S3 simulated patch runtime	30 719.38 ms	1023.98 ms
S4 simulated patch runtime	40 710.02 ms	1357.00 ms

TABLE V
INDICATIVE UPDATE TRANSFER TIMES UNDER MEASURED DAILY VOLUMES.

Case	Transfer time (s)
S2: 100 median tiles (4.72 MB)	0.19–0.40
S2: 1000 median tiles (47.2 MB)	1.91–3.98
S2: mean daily invalidation (126.73 MB)	5.13–10.69
S3: daily SQL patch (5.07 MB)	0.21–0.43

B. Incremental update workload

Table IV summarizes daily-diff update effort, and Figure 2 shows daily dynamics for additions, removals, and speed-tag changes.

S2 improves refresh locality versus S1 (fewer invalidated partitions), while S3 is more efficient than S4 among database-centered variants (1.33× lower mean patch runtime).

To estimate practical mobile transfer costs, payload transmission time is modeled as:

$$T = \frac{B_{\text{payload}} + B_{\text{overhead}}}{R} \quad (3)$$

with $B_{\text{overhead}} = 0.1 \cdot B_{\text{payload}}$. For European 5G conditions, R is bounded by reported ranges in recent network measurements [22], [23]. Under this model, compact SQL patch transfer for S3 remains faster than broad tile-bundle refresh for S2 at equivalent daily coverage.

C. Deployment choice and track relevance

Under equal weighting of query latency and update effort, S3 is the preferred operating point. The architecture provides low-latency on-device lookup, compact single-file deployment, and lower incremental update overhead than S4.

For ITSC positioning, this evidence primarily supports tracks in *Infrastructure and ITS Technologies* and *Traffic Data*

Analytics and Machine Learning, with secondary relevance for *ITS Validation and Verification*.

VI. THREATS TO VALIDITY

The benchmark uses one fixed probe location and heading, and therefore does not represent all geometric or traffic contexts. The update analysis uses one 30-day window and should be extended with longer seasonal coverage. Finally, the current evaluation reports latency and update effort but does not yet include end-to-end driving-assistance effectiveness metrics in fleet-scale trials.

VII. CONCLUSION

This paper evaluated four runtime data architectures for country-scale, offline-capable speed-limit inference on smartphones. Keeping source data and query logic constant isolates architecture effects: global indexes (S1) are not viable for low-latency mobile deployment, tiled packs (S2) improve locality, and embedded spatial databases dominate latency. Among embedded variants, S3 provides the best latency-update tradeoff and is selected as the default architecture for practical deployment.

Future work will extend the benchmark to route-level workloads, multi-country datasets, and longitudinal weekly checkpoints with intervening daily diffs.

VIII. REPRODUCIBILITY

All preprocessing, scenario construction, and benchmark steps are parameterized and versioned in the accompanying research codebase available at <https://github.com/>. The evaluation reported here is tied to the Germany snapshot dated 2026-02-23, with fixed probe location, heading, query modes, and scenario definitions stated in the methodology and evaluation sections.

REFERENCES

- [1] European Parliament and Council of the European Union, “Regulation (eu) 2019/2144 on type-approval requirements for motor vehicles and their trailers, and systems, components and separate technical units intended for such vehicles,” <https://eur-lex.europa.eu/eli/reg/2019/2144/oj>, 2019, accessed: 2026-02-23.
- [2] European Commission, “Commission delegated regulation (eu) 2021/1958 supplementing regulation (eu) 2019/2144 by laying down detailed rules concerning intelligent speed assistance and driver drowsiness and attention warning systems,” https://eur-lex.europa.eu/eli/reg_del/2021/1958/oj, 2021, accessed: 2026-02-23.
- [3] —, “New mandatory vehicle safety features from july 2024,” https://single-market-economy.ec.europa.eu/news/new-mandatory-vehicle-safety-features-2024-07-05_en, 2024, accessed: 2026-02-23.
- [4] Eurostat, “Passenger cars in the eu (statistics explained; data extracted in january 2026),” <https://ec.europa.eu/eurostat/statistics-explained/SEPDF/cache/25886.pdf>, 2026, accessed: 2026-02-23.
- [5] European Automobile Manufacturers’ Association (ACEA), “Car registrations: +0.8% in 2024; battery electric 13.6% market share in december,” <https://www.acea.auto/pc-registrations/new-car-registrations-0-8-in-2024-battery-electric-13-6-market-share-in-december/>, 2025, accessed: 2026-02-23.
- [6] Eurostat, “Internet access in cities and rural areas,” <https://ec.europa.eu/eurostat/web/products-eurostat-news/w/ddn-20241122-1>, 2024, accessed: 2026-02-23.

- [7] M. A. Quddus, W. Y. Ochieng, and R. B. Noland, "Current map-matching algorithms for transport applications: State-of-the art and future research directions," *Transportation Research Part C: Emerging Technologies*, vol. 15, no. 5, pp. 312–328, 2007. [Online]. Available: <https://doi.org/10.1016/j.trc.2007.05.002>
- [8] P. Newson and J. Krumm, "Hidden markov map matching through noise and sparseness," in *Proceedings of the 17th ACM SIGSPATIAL International Conference on Advances in Geographic Information Systems*, 2009, pp. 336–343. [Online]. Available: <https://doi.org/10.1145/1653771.1653818>
- [9] Y. Lou, C. Zhang, Y. Zheng, X. Xie, W. Wang, and Y. Huang, "Map-matching for low-sampling-rate gps trajectories," in *Proceedings of the 17th ACM SIGSPATIAL International Conference on Advances in Geographic Information Systems*, 2009, pp. 352–361. [Online]. Available: <https://doi.org/10.1145/1653771.1653820>
- [10] C. Y. Goh, J. Dauwels, N. Mitrovic, M. T. Asif, A. Oran, and P. Jaillet, "Online map-matching based on hidden markov model for real-time traffic sensing applications," in *2012 15th International IEEE Conference on Intelligent Transportation Systems*, 2012, pp. 776–781. [Online]. Available: <https://doi.org/10.1109/ITSC.2012.6338627>
- [11] Z.-c. He, S. Xi-wei, L.-j. Zhuang, and P.-l. Nie, "On-line map-matching framework for floating car data with low sampling rate in urban road networks," *IET Intelligent Transport Systems*, vol. 7, no. 4, pp. 404–414, 2013. [Online]. Available: <https://doi.org/10.1049/iet-its.2011.0226>
- [12] M. Rahmani and H. N. Koutsopoulos, "Path inference from sparse floating car data for urban networks," *Transportation Research Part C: Emerging Technologies*, vol. 30, pp. 41–54, 2013. [Online]. Available: <https://doi.org/10.1016/j.trc.2013.02.002>
- [13] T. Hunter, P. Abbeel, and A. Bayen, "The path inference filter: Model-based low-latency map matching of probe vehicle data," *IEEE Transactions on Intelligent Transportation Systems*, vol. 15, no. 2, pp. 507–529, 2014. [Online]. Available: <https://doi.org/10.1109/TITS.2013.2282352>
- [14] M. Quddus and S. Washington, "Shortest path and vehicle trajectory aided map-matching for low frequency gps data," *Transportation Research Part C: Emerging Technologies*, vol. 55, pp. 328–339, 2015. [Online]. Available: <https://doi.org/10.1016/j.trc.2015.02.017>
- [15] G. R. Jagadeesh and T. Srikanthan, "Online map-matching of noisy and sparse location data with hidden markov and route choice models," *IEEE Transactions on Intelligent Transportation Systems*, vol. 18, no. 9, pp. 2423–2434, 2017. [Online]. Available: <https://doi.org/10.1109/TITS.2017.2647967>
- [16] M. Che, Y. Wang, C. Zhang, and X. Cao, "An enhanced hidden markov map matching model for floating car data," *Sensors*, vol. 18, no. 6, p. 1758, 2018. [Online]. Available: <https://doi.org/10.3390/s18061758>
- [17] A. Guttman, "R-trees: A dynamic index structure for spatial searching," in *Proceedings of the 1984 ACM SIGMOD International Conference on Management of Data*, 1984, pp. 47–57. [Online]. Available: <https://doi.org/10.1145/602259.602266>
- [18] Geofabrik GmbH, "Geofabrik openstreetmap data extracts – germany latest," <https://download.geofabrik.de/europe/germany-latest.osm.pbf>, 2026, accessed: 2026-02-23.
- [19] OpenStreetMap Wiki Contributors, "Key:highway," <https://wiki.openstreetmap.org/wiki/Key:highway>, 2026, accessed: 2026-02-23.
- [20] —, "Key:maxspeed," <https://wiki.openstreetmap.org/wiki/Key:maxspeed>, 2026, accessed: 2026-02-23.
- [21] Bundesministerium der Justiz, "Straßenverkehrs-ordnung (stvo), section 3 (geschwindigkeit)," https://www.gesetze-im-internet.de/stvo_2013/_3.html, 2013, accessed: 2026-02-23.
- [22] Opensignal, "Europe's mobile network experience 2024," <https://www.opensignal.com/reports/2024/08/europe-mobile-network-experience>, 2024, accessed: 2026-02-23.
- [23] European Commission, "Impact assessment report accompanying the white paper 'how to master europe's digital infrastructure needs?'" (swd(2026) 13 final)," <https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=CELEX:52026SC0013>, 2026, accessed: 2026-02-23.

FIGURES




Fig. 1. Scenario-first latency view. For each scenario (S1–S4), online and offline columns are adjacent; dashed lines delimit scenario groups.




Fig. 2. Daily update dynamics over the 30-day analysis window.